



Web Load Balancing on a Budget

Pain

- Hosting 60+ websites
- Single web server
 - Redundant subsystems (disk, power)
- SPOF
- Inconvenient maintenance windows
 - Clients
 - MY TEAM!


Scope

- Simple. Availability.
 - Minimize/mitigate downtime
 - Outages
 - Planned maintenance
- Session state persistence (for failed web backends) was not required

Stack

- Vmware ESXi (all hosts)
- CentOS (load balancers)
 - Nginx (SSL termination)
 - Haproxy (web load balancing)
 - Keepalived (VRRP)
- Windows Server 2008 (IIS)

Design



```
1 # Let us bind to addresses and ports that may not be "real" (nginx and haproxy need this)
2 net.ipv4.ip_nonlocal_bind = 1
3
4 # Reuse sockets that are in a TIME_WAIT state so we don't exhaust resources
5 net.ipv4.tcp_tw_reuse = 1
6
7 # Give us a larger useable port range (default: 32768 61000)
8 net.ipv4.ip_local_port_range = 1024 65535
9
10 # Close TCP FIN connections faster to help lower resources used by the network stack
11 net.ipv4.tcp_fin_timeout = 30
12
13 # Increase the number of slots that iptables has for tracking connections
14 net.ipv4.netfilter.ip_conntrack_max = 131072
15
16 # Decrease the time that iptables waits to close sockets in TIME_WAIT (def: 120)
17 net.ipv4.netfilter.ip_conntrack_tcp_timeout_time_wait = 30
```

```
1 http {
2     server_tokens off;                # make the hackers work for it
3
4     server_name_in_redirect off;      # use requested Host header
5
6     proxy_read_timeout 1500;          # 25 min; required for long-running reports
7
8     upstream www.example.com {
9         server 192.168.1.2:80;
10    }
11
12    server {
13        listen 192.168.1.2:443;
14        access_log /var/log/nginx/access.log;
15        ssl on; # go!
16        ssl_certificate /usr/local/ssl/wildcard.example.com.cert;
17        ssl_certificate_key /usr/local/ssl/wildcard.example.com.key;
18        ssl_prefer_server_ciphers on;   # prefer SSLv3 and TLSv1 ciphers
19        ssl_ciphers HIGH:+MEDIUM:!ADH:!MD5; # use 128-bit and higher ciphers; exclude ADH, MD5
20        ssl_protocols TLSv1 SSLv3;      # guarantee only TLSv1 and SSLv3 protocols
21
22        location / { # match all HTTP requests
23            proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
24            proxy_set_header secureCookie YES;
25            proxy_pass http://www.example.com;
26        }
27    }
28
29 }
```

```
1 global
2     maxconn 32500
3     nbproc 1
4
5 defaults
6     log global          # log all proxy instances
7     option httplog      # verbose HTTP logging
8     mode http           # layer 7 proxy
9
10    retries 2
11    timeout connect 2s   # 2 retries x 2s timeout = 4s until haproxy looks for new backend
12    option redispatch    # redispatch requests originally sent to downed backend servers
13
14    timeout client 5s
15    timeout server 1500s # 25 minutes, for long-running reports (matches nginx config)
16
17    balance roundrobin   # also leastconn
18
19    option forwardfor    # X-Forward-For header
20    option httpclose     # force HTTP connections closed so all connections are logged
21
22 listen statsWWW 192.168.1.1:8080
23     stats enable        # enable the haproxy stats page (uri-stem = /haproxy?stats)
24     stats auth god:sex
25     stats refresh 120s
26
27 listen www.example.com
28     bind 192.168.1.2:80
29     acl acl_port_80 dst_port eq 80
30     acl acl_secure hdr(secureCookie) YES
31     acl acl_sourceLocal src 192.168.1.2
32     acl acl_excludeCSS url_dir -i CSS
33     redirect location https://www.example.com/sslRedirect.asp code 301 if acl_port_80 !acl_secure !
acl_sourceLocal !acl_excludeCSS
34     cookie SERVERID insert indirect nocache
35     server www1 192.168.1.3 cookie www1 weight 1 check inter 1s fall 3 rise 2
35     server www2 192.168.1.4 cookie www2 weight 1 check inter 1s fall 3 rise 2
35     server www3 192.168.1.5 cookie www3 weight 1 check inter 1s fall 3 rise 2
38     option httpchk GET /ping.htm HTTP/1.1\r\nHost:www.example.com    # health check
39     capture request header X-Forwarded-For len 15
```



```
1 global_defs {
2     lvs_id lb1
3 }
4
5 vrrp_sync_group virtualGroup1 {
6     group {
7         virtualInstance1
8     }
9 }
10
11 vrrp_instance virtualInstance1 {
12     interface eth0
13     advert_int 1          # advertise interval in seconds
14     state MASTER          # either MASTER or BACKUP
15     virtual_router_id 1   # MUST be between 1 and 255
16     priority 101          # 101 on MASTER, 100 on BACKUP
17     authentication {
18         auth_type AH
19         auth_pass "sexdrugsandrokickandroll"
20     }
21     virtual_ipaddress { # only 20 IPs per block, blah
22         192.168.1.1/24 brd 192.168.1.255 dev eth0
23         192.168.1.2/24 brd 192.168.1.255 dev eth0
24         192.168.1.3/24 brd 192.168.1.255 dev eth0
25         192.168.1.4/24 brd 192.168.1.255 dev eth0
26         ...
27         192.168.1.19/24 brd 192.168.1.255 dev eth0
28     }
29 }
```

Pitfalls




- VMware and TSO
 - When jumbo frames go bad
 - HTTP 500 errors
 - Web servers “offline” for 20-30 minutes
 - Disable TSO
 - <http://www.ryanfrantz.com/2011/02/03/tcp-segmentation-offload/>


What's Happening?

- No visibility into performance
 - Request volume (nginx/haproxy/IIS)
 - Seasonality
- Health (beyond up/down)
 - keepalived

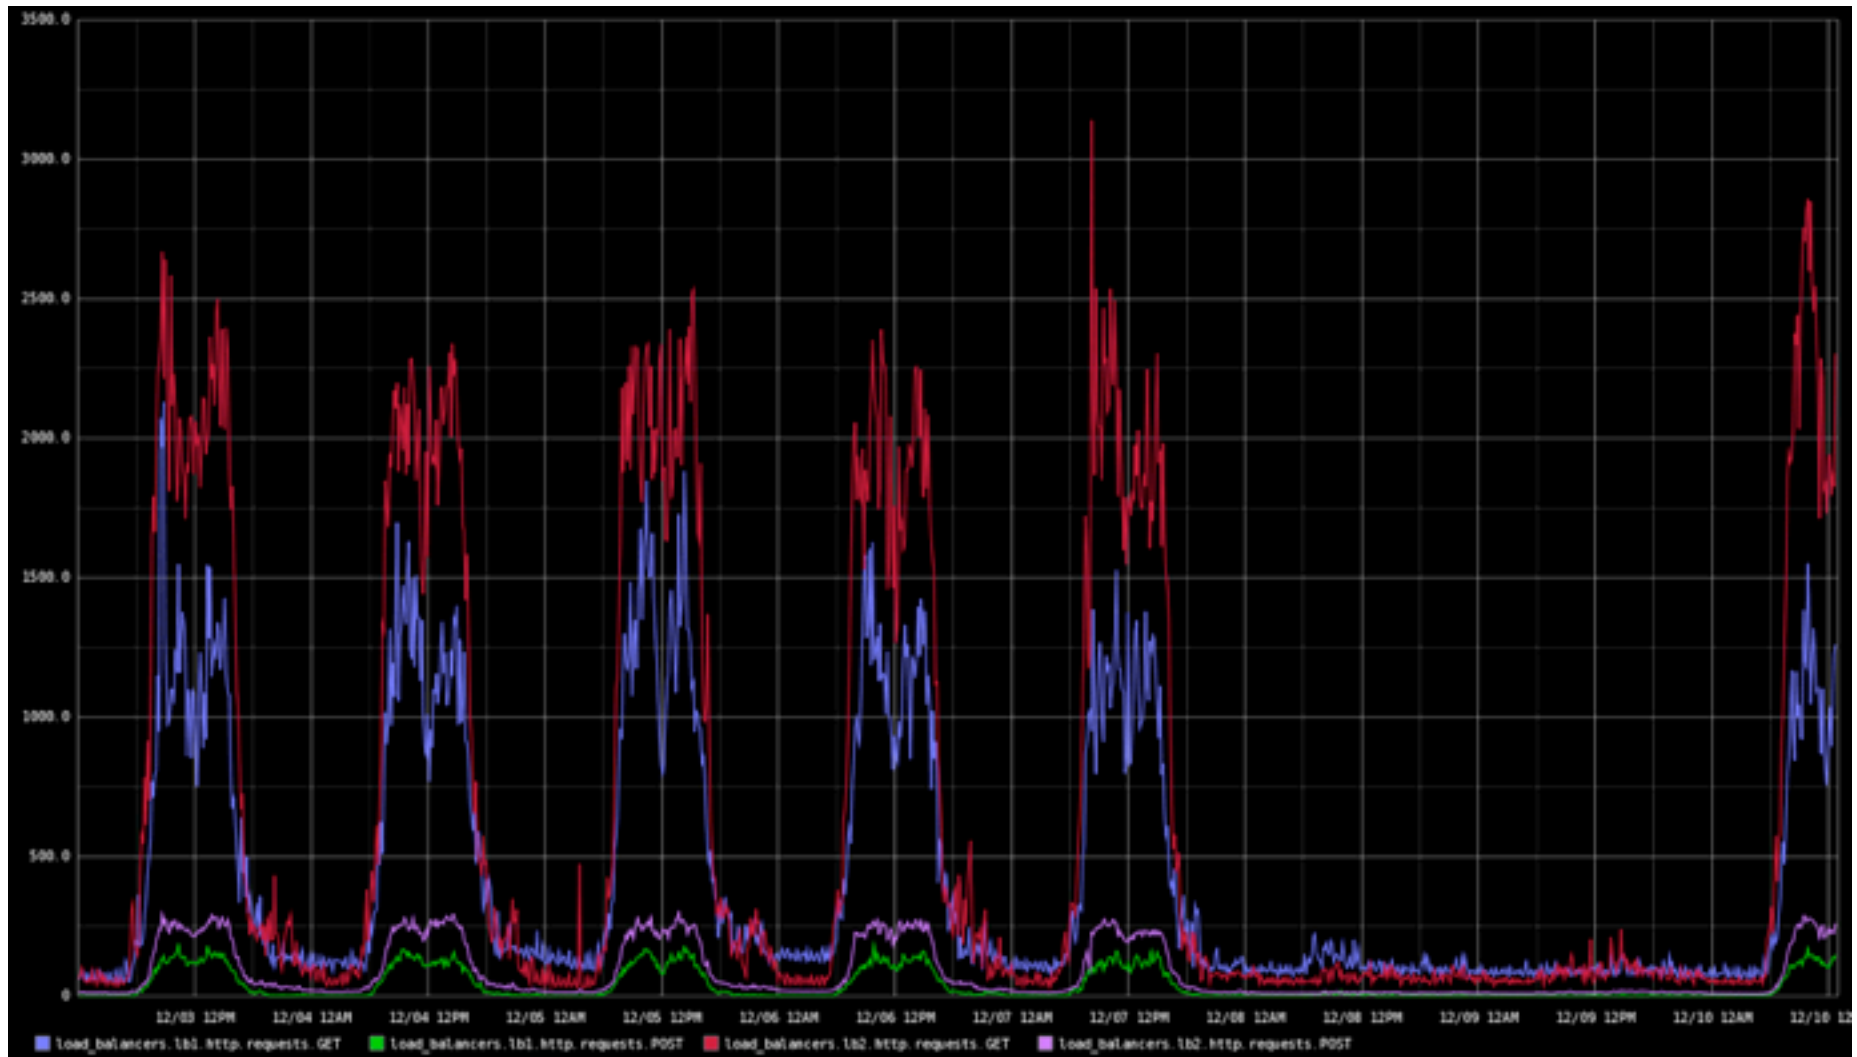
HTTP Traffic: One Load Balancer



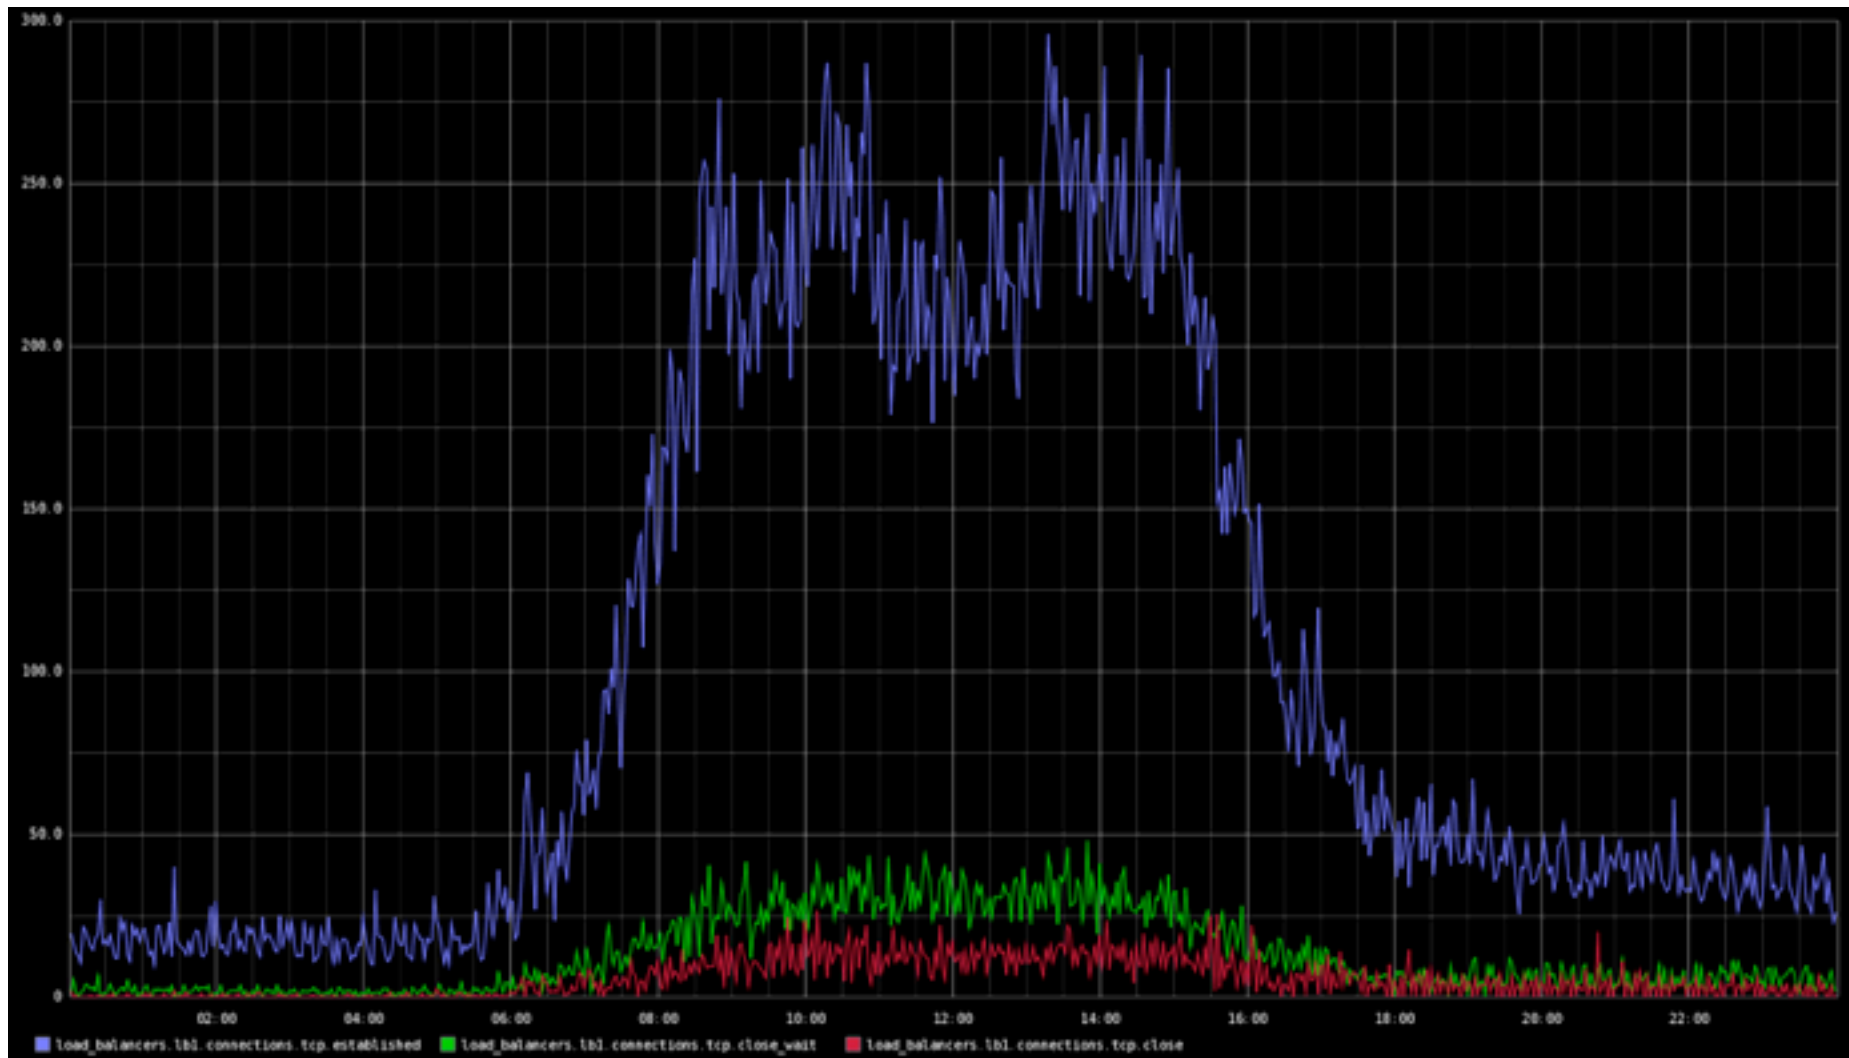
HTTP Traffic: Two Load Balancers



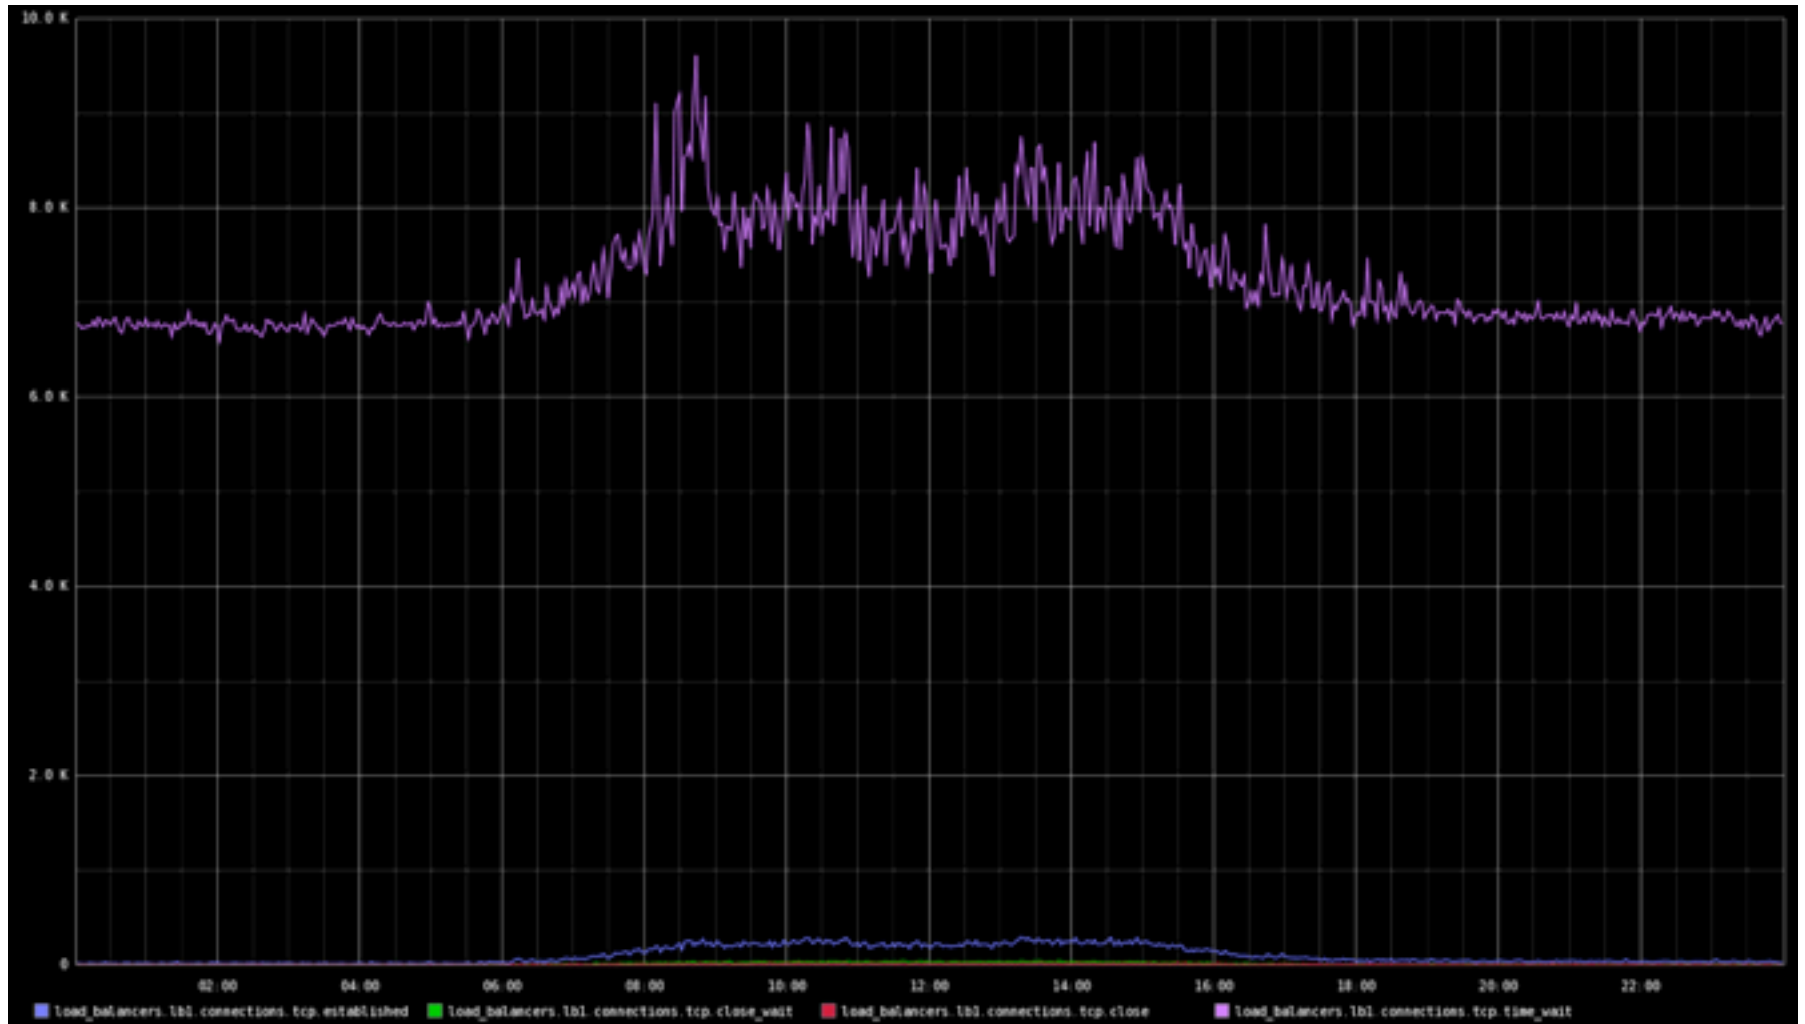
HTTP Traffic: One Week



TCP Connections



TCP Connections: + TIME_WAIT



Thank You for Attending LOPSA–East '13

Please fill out the Trainer Evaluation
<http://lopsa-east.org/2013/training-survey>

Rate LOPSA–East '13
<http://www.lopsa-east.org/2013/rate-lopsa-east-13>

[@Ryan_Frantz](#)
rfrantz@etsy.com
github.com/RyanFrantz
www.ryanfrantz.com

